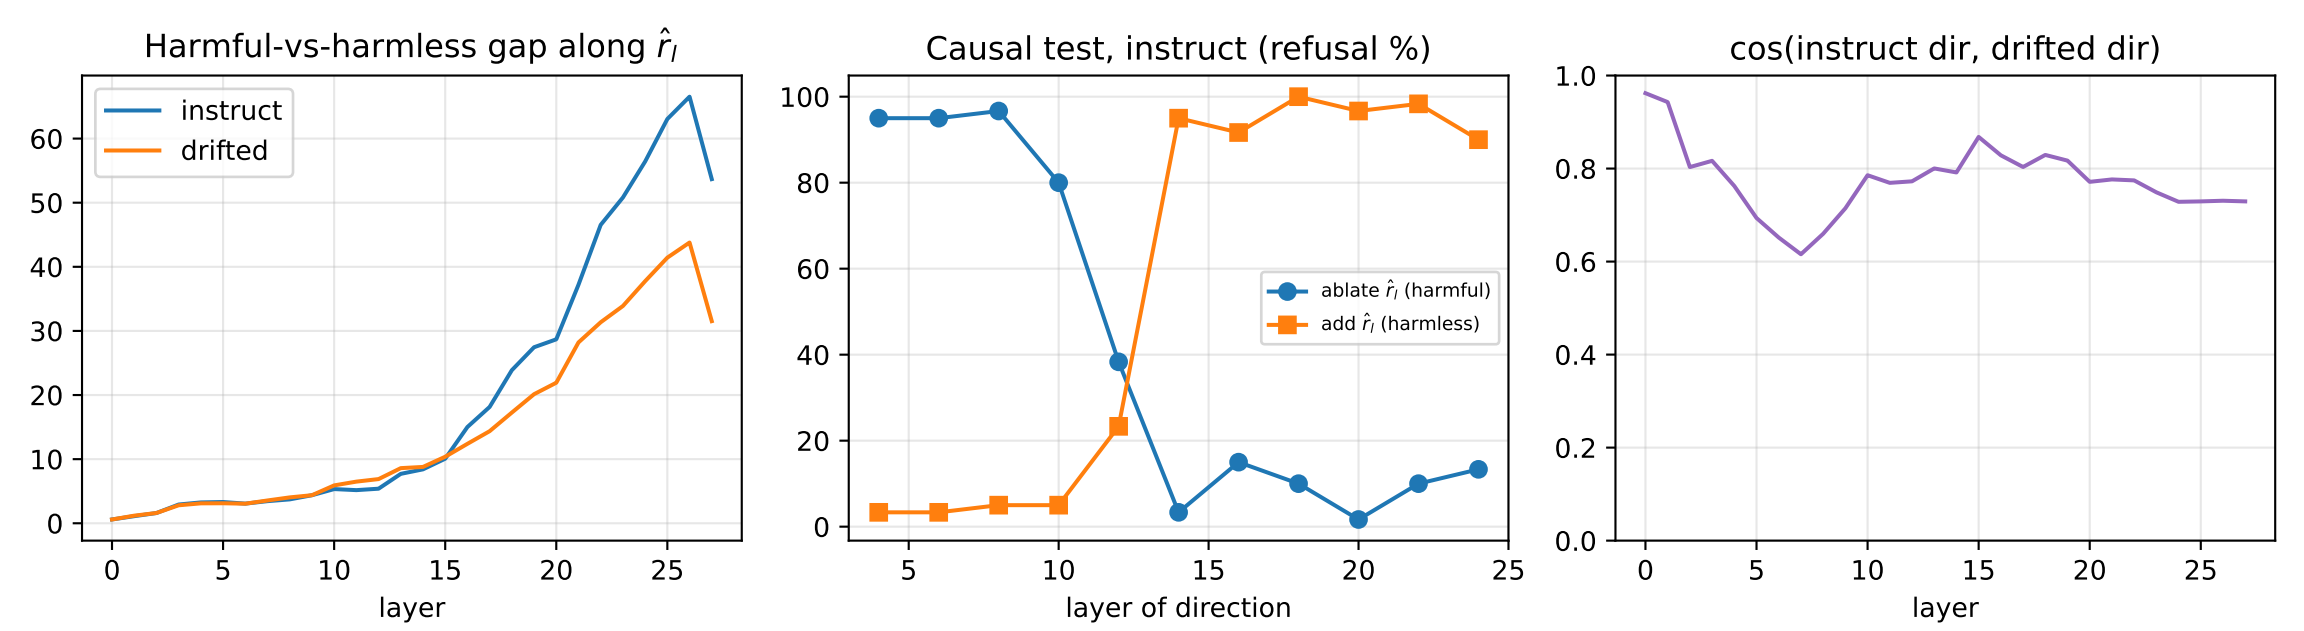
Harmful-vs-harmless gap along $\hat{r}_l$
instruct
drifted
layer
Causal test, instruct (refusal %)
ablate $\hat{r}_l$ (harmful)
add $\hat{r}_l$ (harmless)
layer of direction
cos(instruct dir, drifted dir)
layer