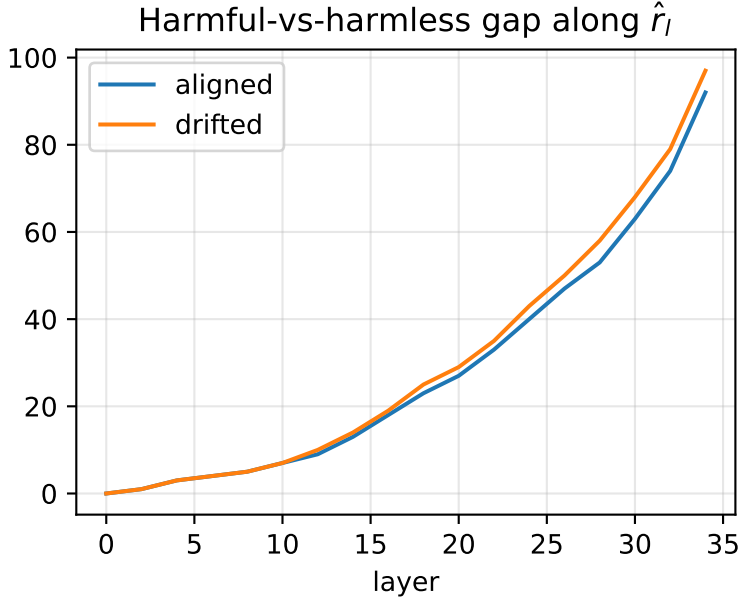
Harmful-vs-harmless gap along $\hat{r}_l$
aligned
drifted
layer

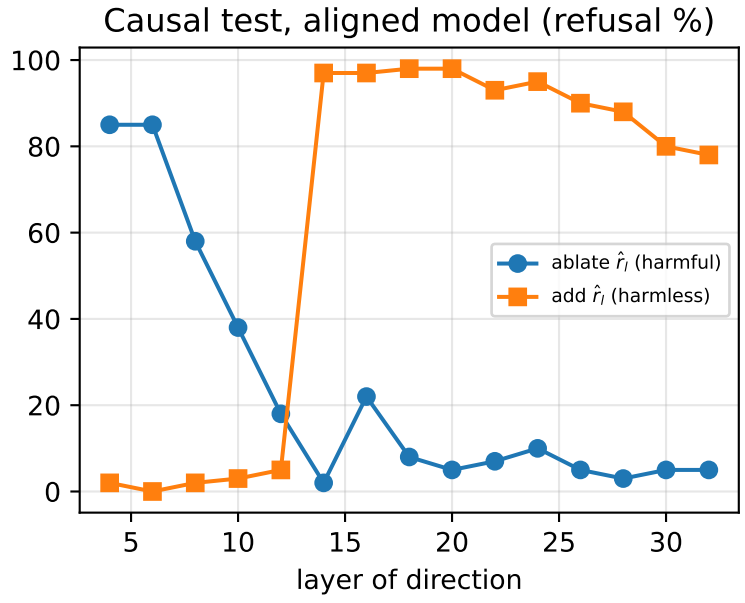
Causal test, aligned model (refusal %)
ablate $\hat{r}_l$ (harmful)
add $\hat{r}_l$ (harmless)
layer of direction

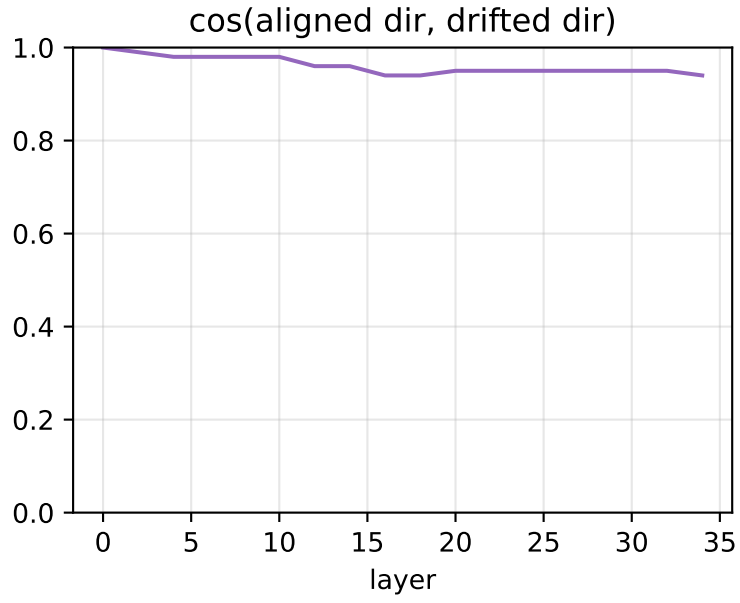
cos(aligned dir, drifted dir)
layer